

Docker From The Ground Up

Part 3

Multihost, Swarm and
Production





Thanks to our Sponsor!

BlackCat /





Who's this guy?

Matt Todd

Principal Architect @ Resonate.tech

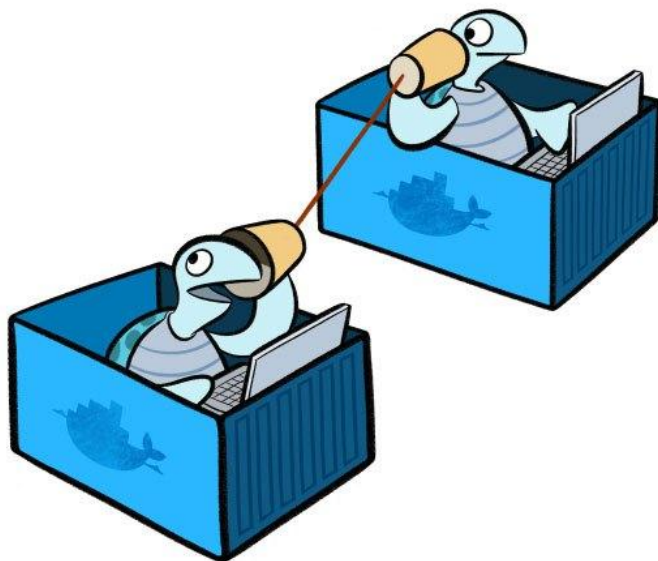
Cloud Native Advocate

`github.com/mattjtodd`

`hub.docker.com/u/mattjtodd`



Recap!





Containers

- Linux Namespaces and CGroups
- Process resource isolation and constraint definition
- Share host CPU / Memory
- Lighter weight than VMs
- Fast startup times
- Docker Portability / Abstraction



Networking

- Namespaced networking stack per container (can be shared)
- Isolated by design
- Extensive use of bridge networking (Virtual switches)
- Multi-host capability (with swarm)
- Portable and able to work in many networking stacks
- Network Types (None, Host, Bridge, User-defined, MacVlan, Overlay)



Docker Compose

- A way of defining and managing multiple containers and their resources
- Use YAML templates to define multiple services including networks, volumes
- Start and develop applications using an iterative lifecycle
- Namespaced resources provides services stack isolation
- Single host but can be used with Swarm
- Installed with Mac / Windows, separately on Linux

Into
Production!



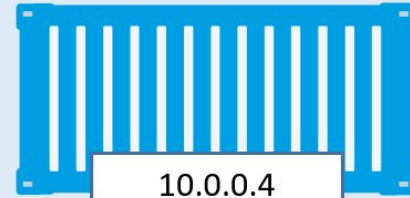
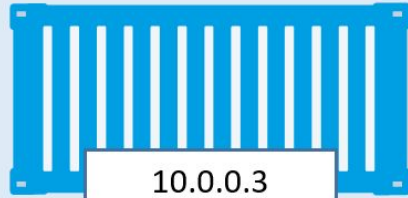


Overlay Networks

- A distributed network between hosts
- Requires swarm mode to be enabled for participant nodes
- VXLANs to create a virtual Layer 2 network on top of an existing Layer net
- Can be encrypted via IPSEC tunnels via driver config
- Creates and *ingres* (default) and *docker_gwbridge* networks

node1

node2



Overlay 10.0.0.0/24

172.31.1.5

192.168.1.25

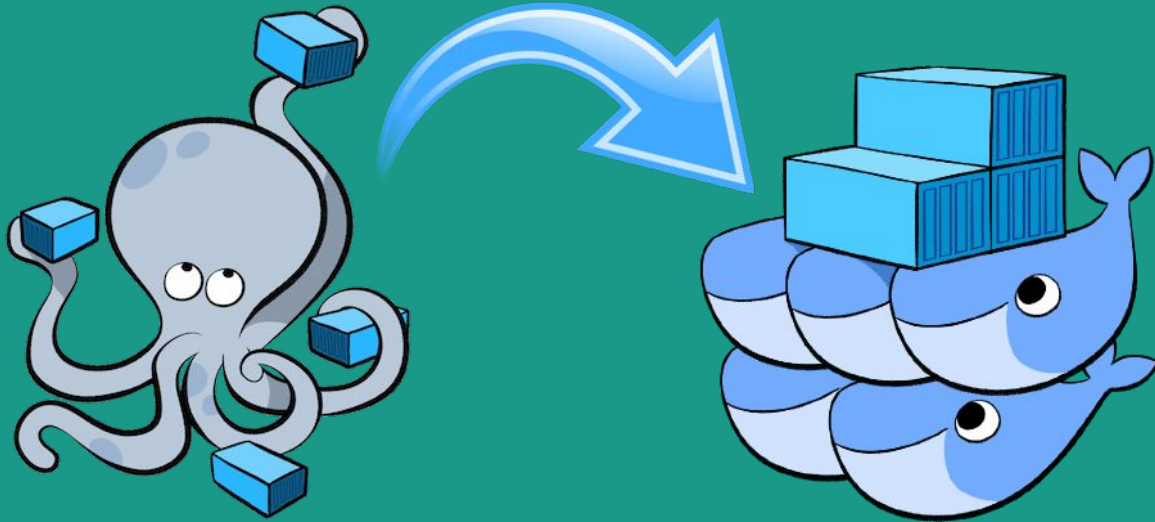




Tips!

- Check for IP range conflicts in *ingres* and *docker_gwbridge*
- Firewall configured:
 - TCP port 2377 for cluster management communications
 - TCP and UDP port 7946 for communication among nodes
 - UDP port 4789 for overlay network traffic
- Always test performance of your network!
 - IPSec / Kernel issues!
- Swarm initialized? :)

Docker Swarm



<https://github.com/mattjtodd>

docker-from-the-ground-up-part-3



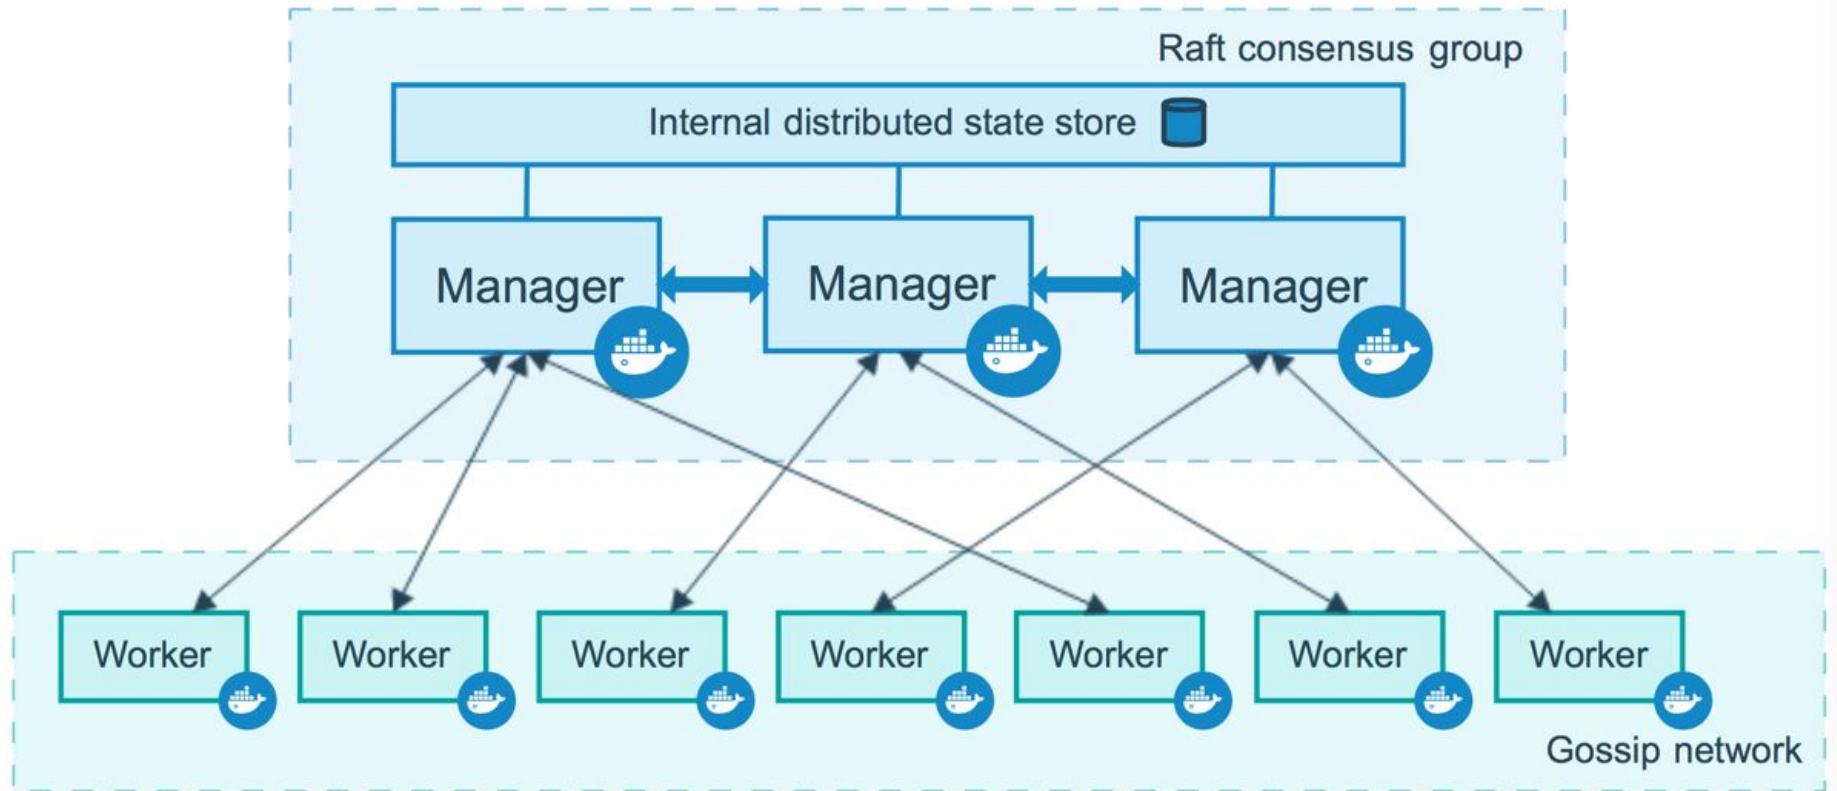
Desirable Characteristics

- Resilient to failure
- Secure in the cloud
- Horizontally scalable
- Easily managed
- Operator Friendly
- Heterogeneous resource capability
- Service Discovery



Features

- Built into Engine; batteries included!
- Orchestration via cli driven services / compose stacks
- Multi-host service mesh routing / load balancing
- DNS based Service discovery
- Config / Secrets management
- Rolling update / rollback / spread policies





Swarm

- A cluster will be at least one node (preferably more)
- A node can be a manager or a worker
- A manager actively takes part in the Raft consensus, and keeps the Raft log
- One manager is elected as the leader; other managers merely forward requests to it
- The workers get their instructions from the managers
- Both workers and managers can run containers



Initialisation

- Initialize a seed manager node
- Join other nodes with the join tokens
- Specify interface ip if needed
- Vagrantfile in **vagrant/docker**



Initialisation

```
$ vagrant up
$ vagrant ssh node-1
$ docker swarm init --advertise-addr=192.168.50.150
$ docker node ls
Copy the join token command and exit the ssh session
$ vagrant ssh node-2
$ <paste token join command in here>
$ docker node ls # we're a worker, so can't see this
$ exit
$ vagrant ssh node-2
$ docker node promote
```



Orchestration

- Manager(s) job:
 - Schedule and maintain containers on a cluster of workers
- Worker(s) job:
 - Execute containers and report back on the status
- Services define the tasks to be run
 - Image, port mappings replicas, ...
- Tasks are the encapsulation for a container
 - Assigned to a node and lives there until failure



Services & Tasks

- Services define the *shape* and *characteristics* of tasks
- Health-checks define when services should be registered in the task pool
- Placement restrictions
- Resource constraints
- Rolling update policies
- Rollback policies
- Modes *replicated* and *global*

Docker Engine client

docker service create

swarm manager

RAFT

API

accepts command and creates service object

orchestrator

reconciliation loop that creates tasks for service objects

allocator

allocates ip addresses to tasks

dispatcher

assigns tasks to nodes

scheduler

instructs a worker to run a task

worker node

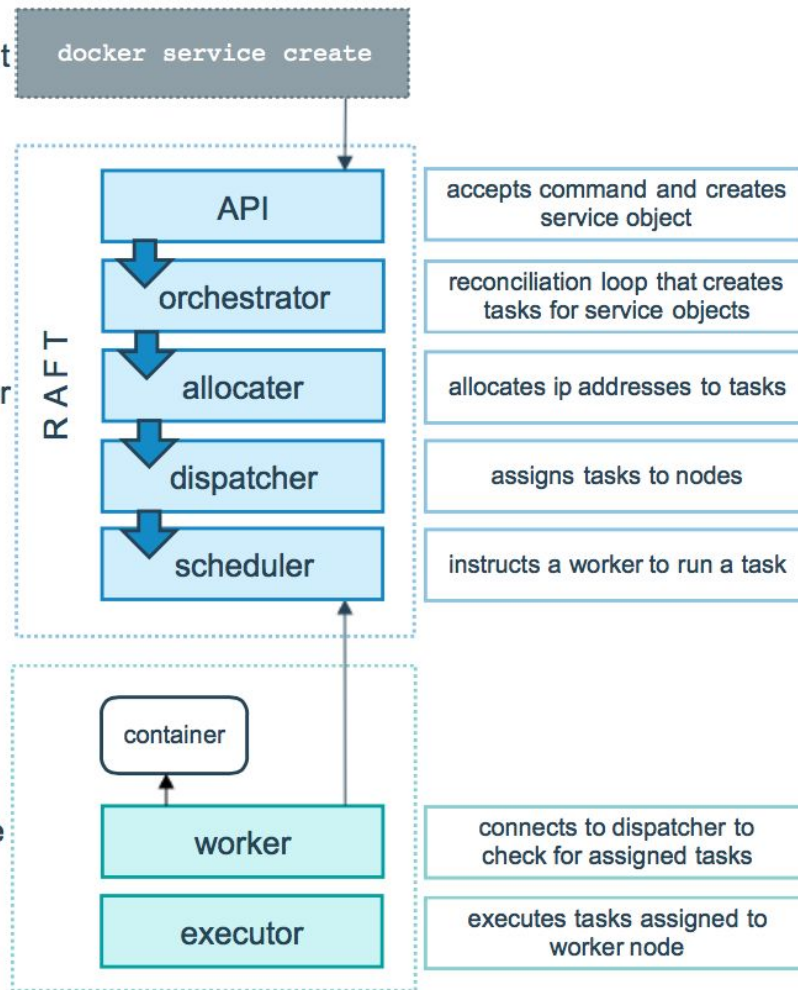
container

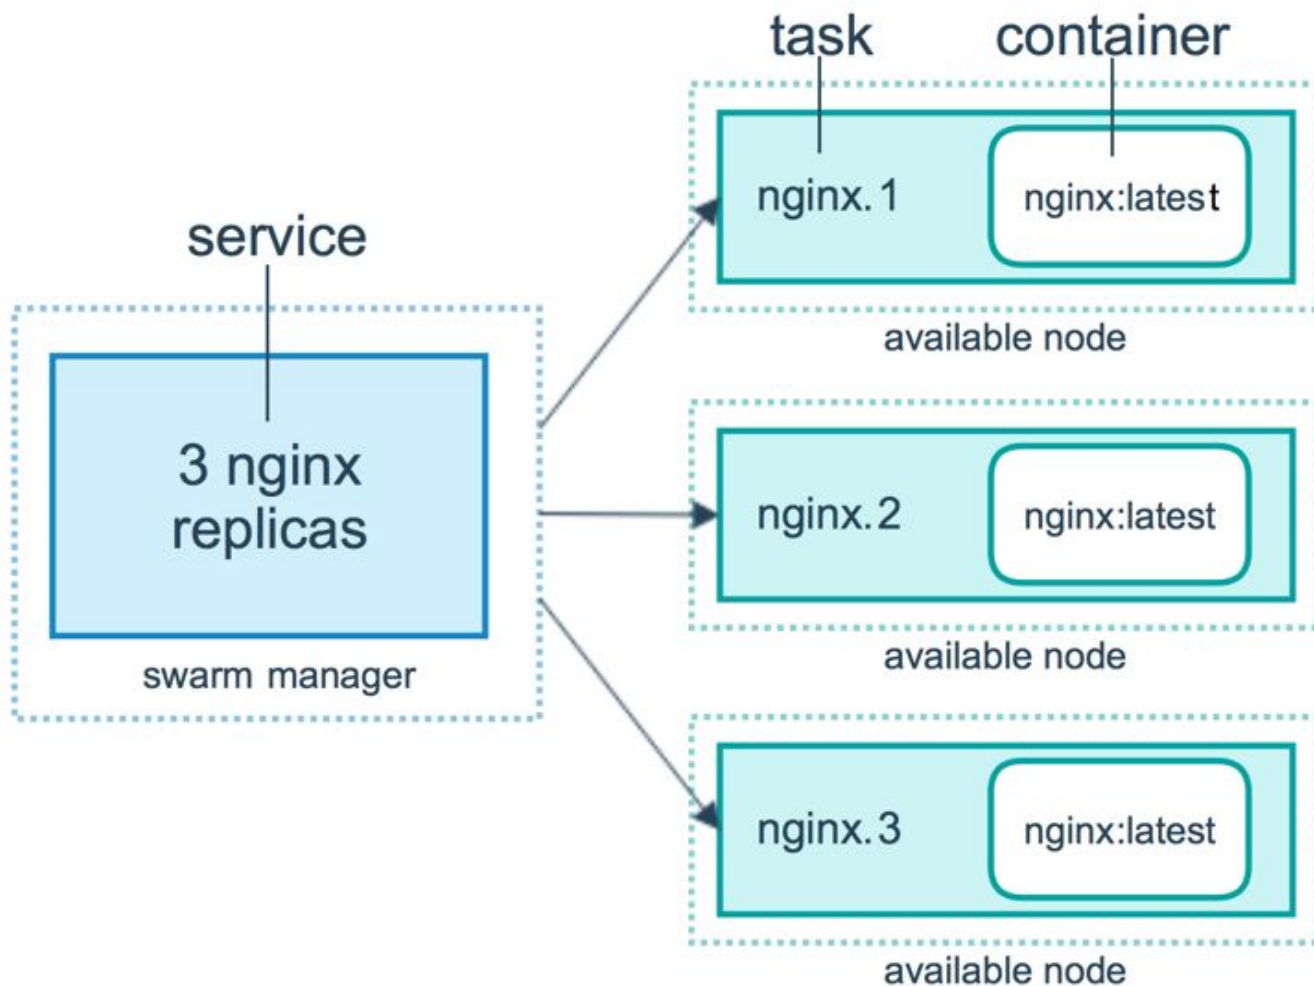
worker

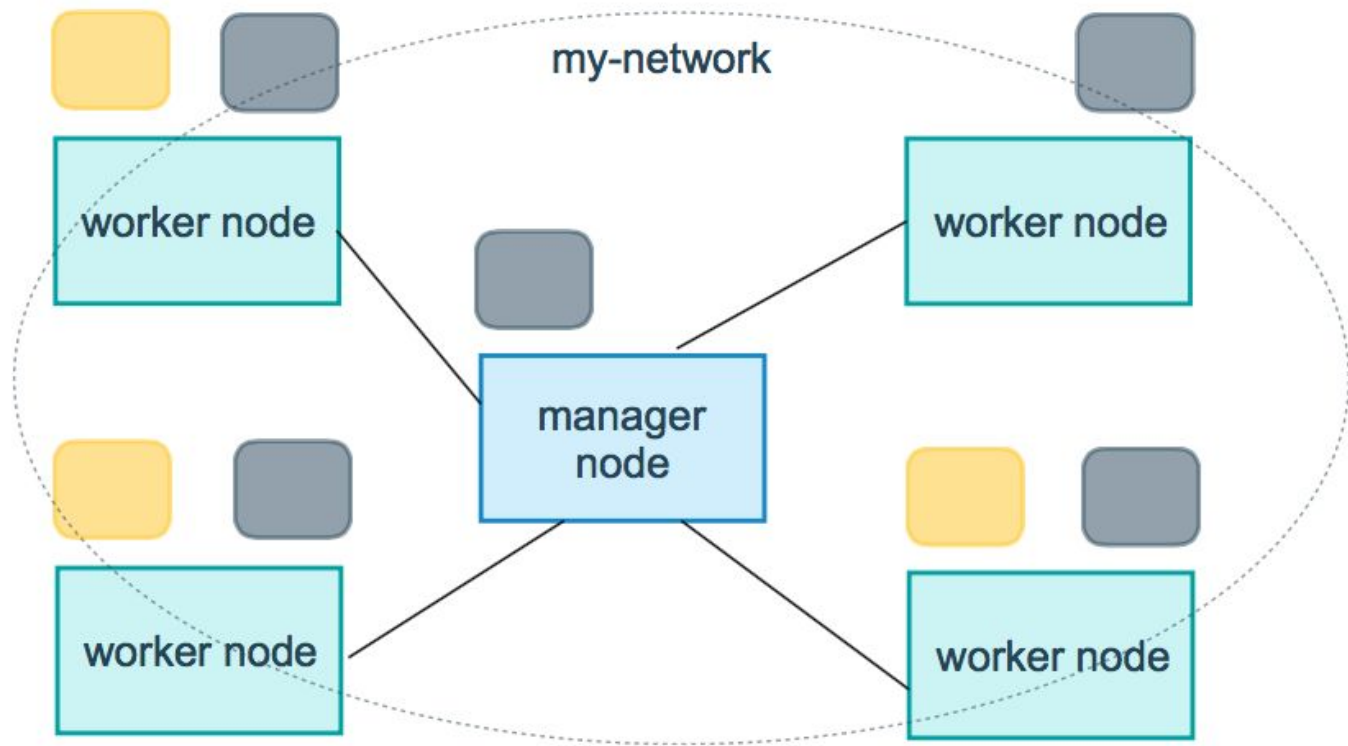
connects to dispatcher to check for assigned tasks

executor

executes tasks assigned to worker node







replicated service
with 3 replicas



global service with
replicas on every node



Service Discovery

- As with compose Docker uses DNS to register services in a pool
- Only *healthy* tasks are registered in the pool
- User Defined bridge networks register named containers here
- Overlay has Two types:
 - VIP (default)
 - DNSRR



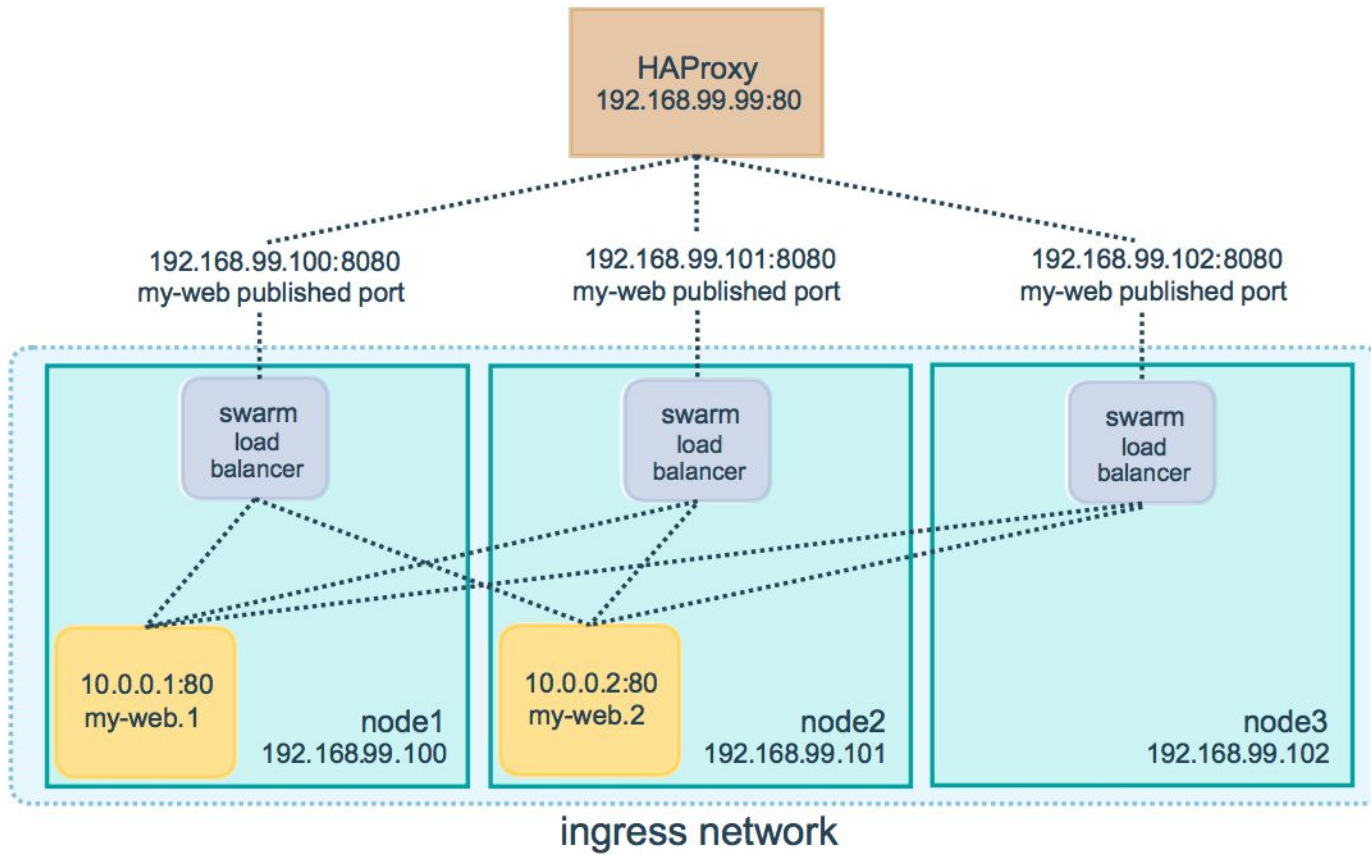
Service Discovery

- VIP
 - Resolves to a single VIP which load balances between tasks
 - Can be part of the ingres service mesh
 - Pool DNS lookup takes the form of `tasks.<service-name>`
- DNSRR
 - Cannot be part of the service mesh
 - Round-robin ip of containers
 - Self configured LB based on DNS A query
 - Pool DNS lookup by `<service-name>` directly



Load Balancing / Service Mesh

- Containers can live on any suitable host
- Physical / Virtual LB balances load between nodes
- Ingres ports are mapped to each host
- Traffic routed to service tasks in the pool
- Load balance you nodes!



Routing Mesh



Healthcheck

```
healthcheck:  
  test: ["CMD", "curl", "-f", "http://localhost"]  
  interval: 1m30s  
  timeout: 10s  
  retries: 3  
  start_period: 40s
```



Deploy Policies

- Restart Policy
 - If and how to restart on exit
- Rollback Config
 - What to do in event of service update failure
- Update Config
 - How a service should be updated
 - E.g. stop /start first, delay etc.



Secrets & Configs

- Persisted by the managers
- Allocated to containers on as needed basis
- Prevents proliferation of sensitive material
- Keep you containers vanilla!
- Mounted in the containers files system
- Only secrets encrypted in managers



Logging Drivers

- Log aggregation remote / local
- Multiple protocols
 - json-file (default)
 - awslogs
 - fluentd
 - gcplogs
 - gelf
 - Splunk



Placements & Preferences

```
deploy:
  placement:
    constraints:
      - node.role == manager
      - engine.labels.operatingsystem == ubuntu 18.04
    preferences:
      - spread: node.labels.az
```



Stackfiles

- Docker compose files version 3+
- Manage service definitions much like compose
- Deploy directive provides service and task properties
- <https://docs.docker.com/compose/compose-file/>



Sample Stackfile

```
version: '3'
```

```
services:
```

```
  portainer:
```

```
    image: portainer/portainer
```

```
    command: -H "tcp://tasks.portainer-agent:9001" --tlsskipverify --no-auth
```

```
    ports:
```

```
      - "9000:9000"
```

```
    volumes:
```

```
      - /var/run/docker.sock:/var/run/docker.sock
```

```
    deploy:
```

```
      mode: replicated
```

```
      replicas: 1
```

```
      placement:
```

```
        constraints: [node.role == worker]
```



Production Tips

- Use logging aggregation
- Monitor *all* hosts
- Label and tag everything using a consistent schema
- Keep the managers free from work (only control plane agents)

Demo Lab





Goals

- Stand up a 1 manager 3 worker swarm cluster
- Use Portainer to visualise state
- Deploy EFK stack
- Observe Grafana
- Deploy a small application
- Run some test loads
- Scale in / out



Tools

- Vagrant
- Ubuntu Bionic Beaver
- Docker Swarm Mode
- Portainer
- SSH



DOCKER SWARM
Manager
192.168.10.1



Worker 1
192.168.10.2



Worker 2
192.168.10.3



Worker 3
192.168.10.4

Vagrant

Orchestration
SSH Tunnel
SSH Port Forward

Docker CLI over SSH
Forwarded Port

Browser UI over
forwarded ports

Manager

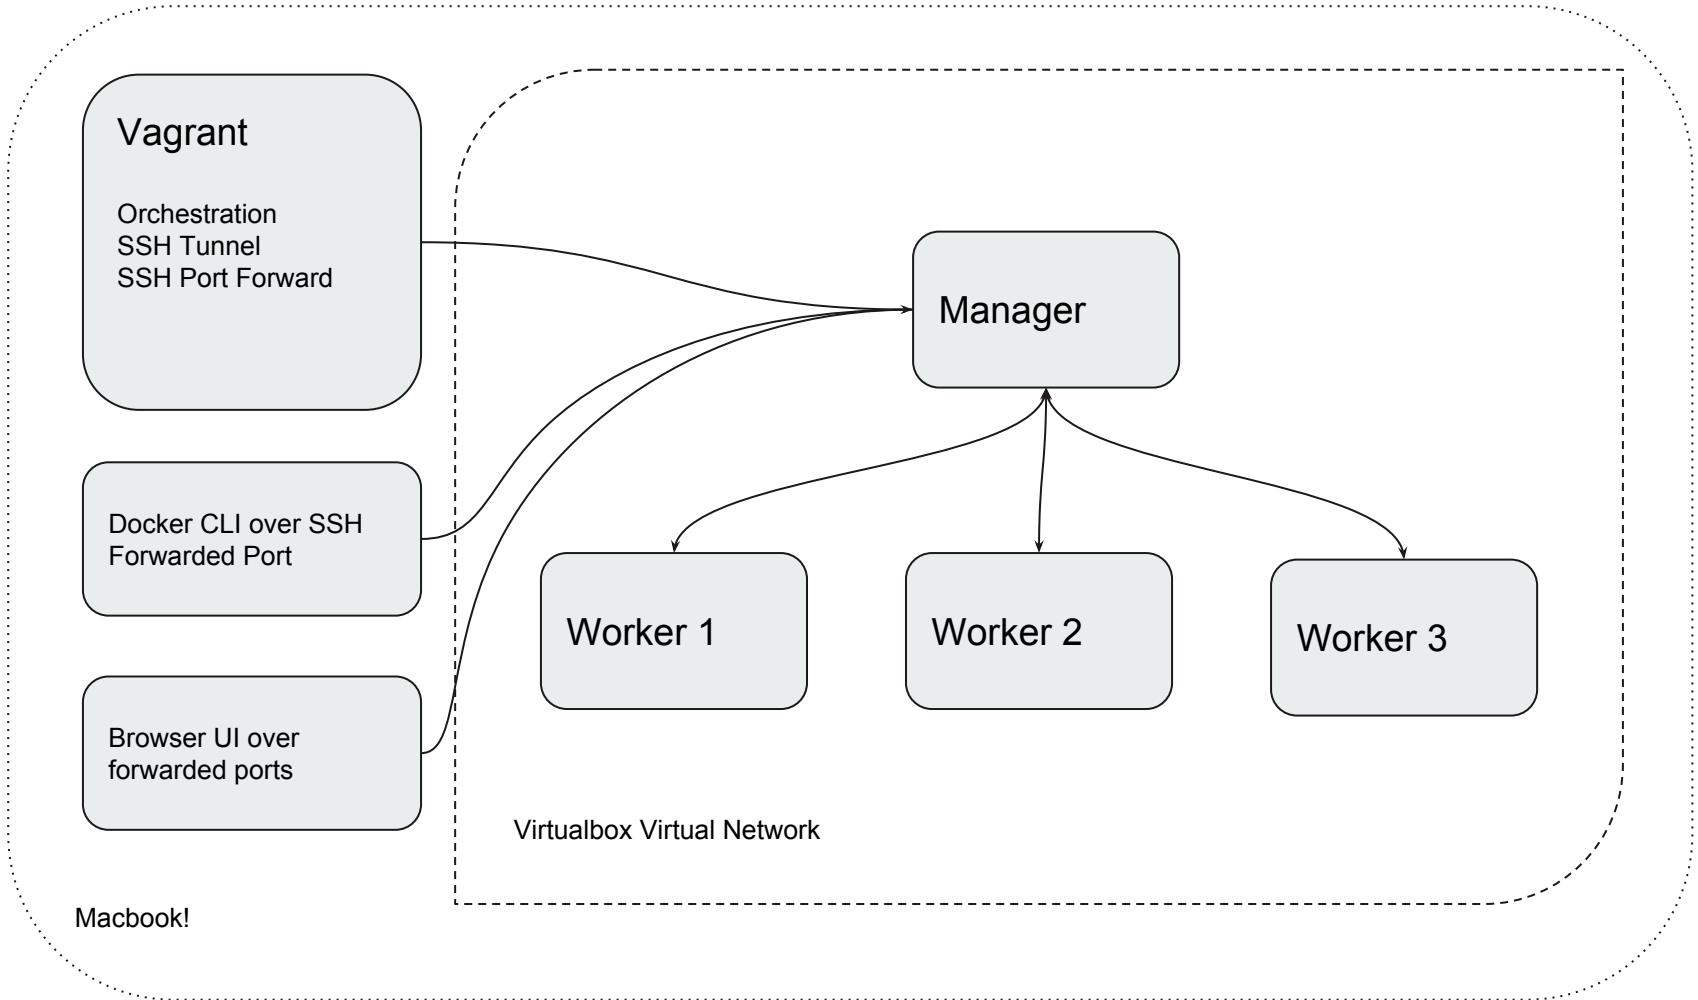
Worker 1

Worker 2

Worker 3

Virtualbox Virtual Network

Macbook!



Manager

Portainer Agent

Worker

Fluentd

Kibana

RabbitMQ Node

Elasticsearch Node

Metricbeat Agent

Portainer Agent

Worker

Spring Client

Grafana

RabbitMQ Node

Elasticsearch Node

Metricbeat Agent

Portainer Agent

Worker

Python Client

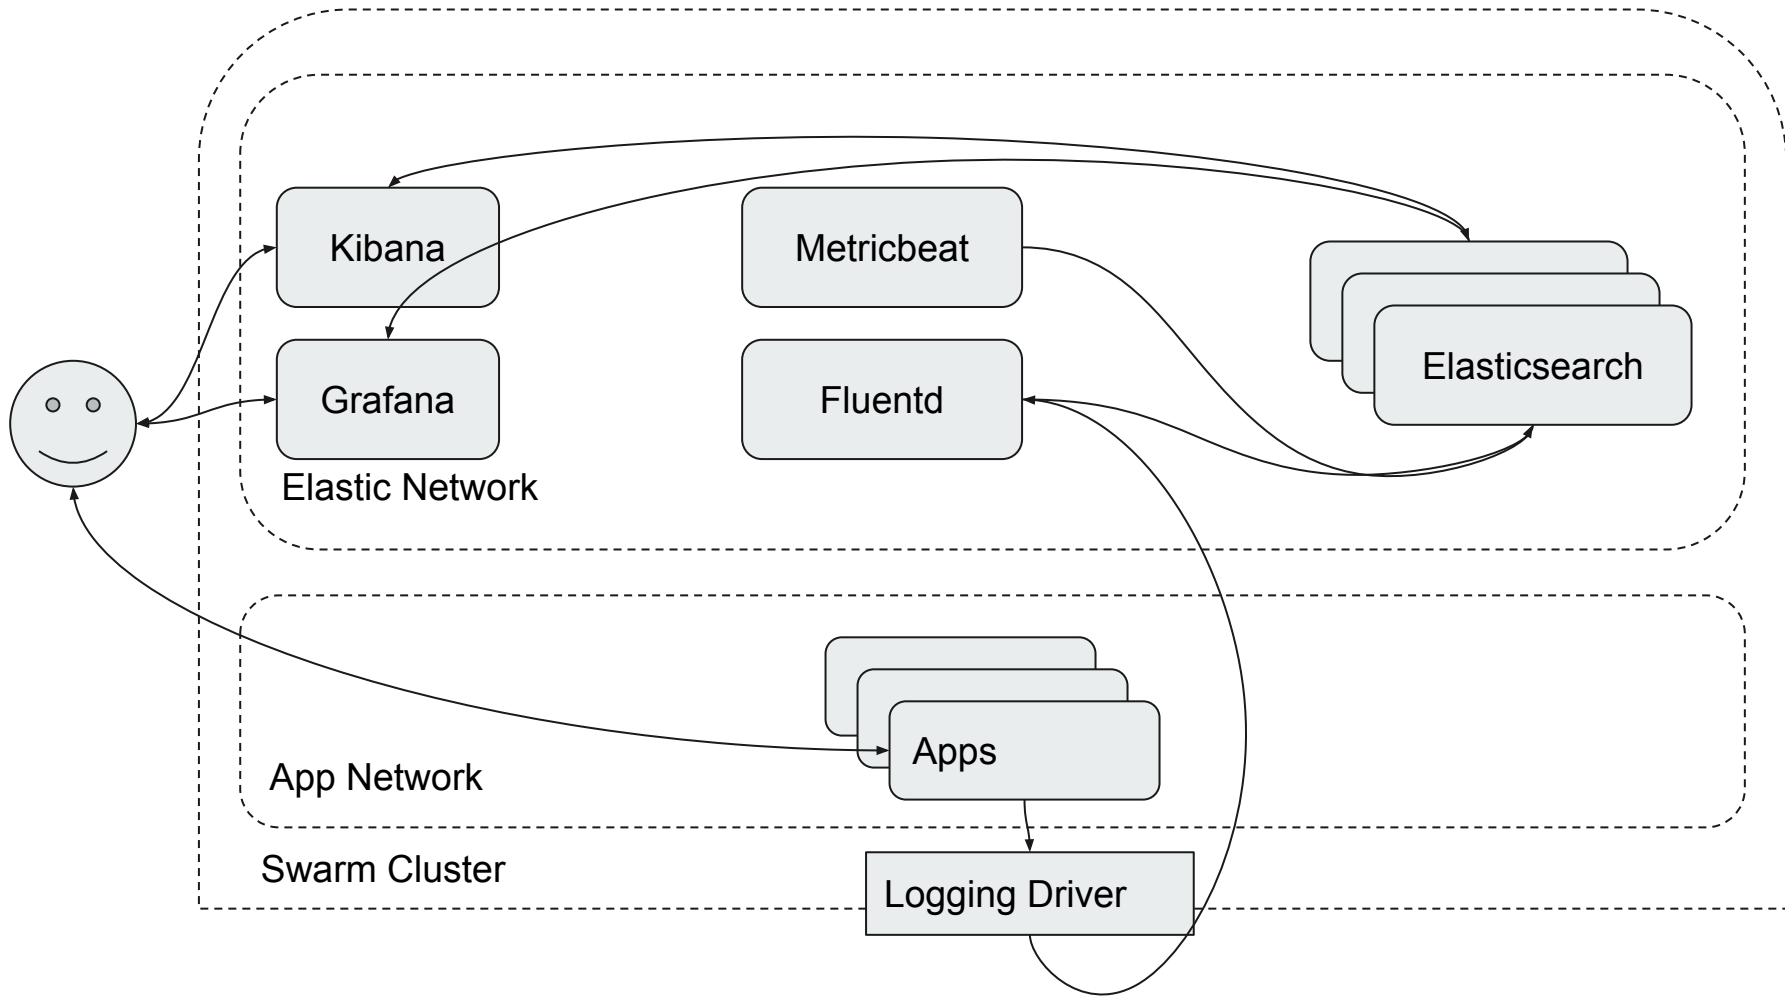
HAProxy (Rabbit)

RabbitMQ Node

Elasticsearch Node

Metricbeat Agent

Portainer Agent





SSH-Fu

- Docker 18.09.0 supports connection over SSH tunnels
 - <https://github.com/docker/cli/pull/1014>
- SSH Port-forwarding will achieve the same
- SSHuttle may also work and resolve hostnames



Getting Started

```
vagrant ssh -- \  
  -L9999:/var/run/docker.sock \  
  -L9000:localhost:9000 \  
  -L5601:localhost:5601 \  
  -L3000:localhost:3000 \  
  -L15672:localhost:15672 \  
  -L5000:localhost:5000 \  
  -L8080:localhost:8080
```



Getting Started

```
# Start up the vagrant managed VMS.  This will create the cluster
$ vagrant up

# Create the SSH Port Forwards for access to the published services
$ vagrant ssh -- \
  -L9999:/var/run/docker.sock \
  -L9000:localhost:9000 \
  -L5601:localhost:5601 \
  -L3000:localhost:3000 \
  -L15672:localhost:15672 \
  -L5000:localhost:5000 \
  -L8080:localhost:8080
```



Deploy EFK

Configure the host ENV variable for the docker client:

```
$ export DOCKER_HOST=localhost:9999
```

Listing the processes via the CLI should yield:

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
PORTS	NAMES			
efe073030701	portainer/agent:latest	"/agent"	10 minutes ago	Up 10
minutes				
portainer_portainer-agent.3ojifkjtru85qe3kcx2zzday7.icmyg91ylw5rfgfze9wztwnn0				
983231ecb5ce	registry:2	"/entrypoint.sh /etc..."	10 minutes ago	Up 10
minutes	5000/tcp	registry-mirror_registry-mirror.1.lkfo7pljrrazsyvkvx84qa4s4p		



Deploy EFK

Move into the **elastic-cluster** dir and run:

```
$ docker stack deploy -c docker-compose.yml efk
```



Deploy EFK

```
$ docker service ls
```

ID	NAME	MODE	REPLICAS
IMAGE	PORTS		
sfwiht8446ch	efk_elasticsearch	global	0/3
elastic/elasticsearch:6.4.3			
d4bgkdj8knz5	efk_fluentd	replicated	0/1
dockerbirmingham/fluentd:latest	*:24224->24224/tcp		
es01tx3x4rjv	efk_grafana	replicated	0/1
grafana/grafana:latest	*:3000->3000/tcp		
r9w8skclnfor	efk_kibana	replicated	0/1
elastic/kibana:6.4.3	*:5601->5601/tcp		
mlhv7iralzbm	efk_metricbeat-host	global	0/4
elastic/metricbeat:6.4.3			
tikymzswwhywa	portainer_portainer	replicated	1/1
portainer/portainer:latest	*:9000->9000/tcp		
f5215am43p98	portainer_portainer-agent	global	4/4
portainer/agent:latest			
vpqnfccch3d95	registry-mirror_registry-mirror	replicated	1/1
registry:2	*:5555->5000/tcp		



Deploy EFK

- Portainer UI
 - <http://localhost:9000>
- Kibana:
 - <http://localhost:5601>
- Grafana:
 - <http://localhost:3000>



Deploy App

Back up to project root and into the **rabbitmq** dir.

```
$ docker stack deploy -c docker-compose.yml app
```

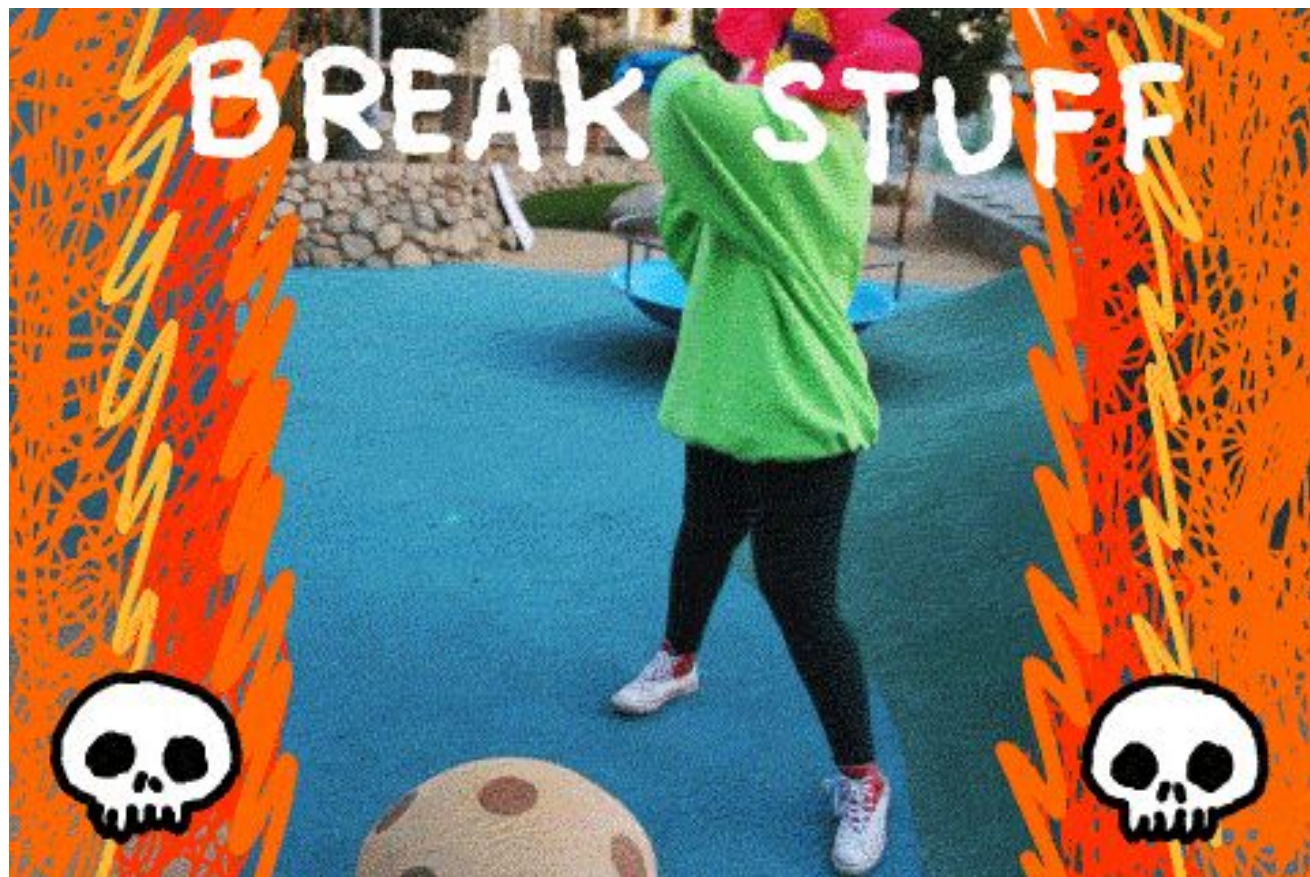
```
$ docker service ls
```

...



Access App

- RabbitMQ Manager
 - Localhost:15672
 - monitor
 - Password
- Python HTTP client
 - <http://localhost:5000/greeting?message=Hello>
- Spring Boot HTTP Client:
 - <http://localhost:8080/greeting?message=Hello>





Additional Resources

<https://www.katacoda.com/courses/docker-orchestration>

<https://container.training/swarm-selfpaced.yml.html>

<https://training.play-with-docker.com/>

<https://docs.docker.com/engine/swarm/>